



(12) **United States Patent**
Shwartsman et al.

(10) **Patent No.:** **US 9,405,545 B2**
(45) **Date of Patent:** **Aug. 2, 2016**

(54) **METHOD AND APPARATUS FOR CUTTING SENIOR STORE LATENCY USING STORE PREFETCHING**

2212/6028; G06F 9/383; G06F 9/3834; G06F 9/3861; G06F 2212/12
See application file for complete search history.

(75) Inventors: **Stanislav Shwartsman**, Haifa (IL);
Melih Ozgul, Hillsboro, OR (US);
Sebastien Hily, Portland, OR (US);
Shlomo Raikin, Ofer (IL); **Raanan Sade**, Kibutz Sarid (IL); **Ron Shalev**, Ceaseria (IL)

(56) **References Cited**

U.S. PATENT DOCUMENTS

5,664,137 A * 9/1997 Abramson G06F 9/30043
712/216
5,996,071 A * 11/1999 White G06F 9/3812
712/207

(Continued)

FOREIGN PATENT DOCUMENTS

(73) Assignee: **Intel Corporation**, Santa Clara, CA (US)

(*) Notice: Subject to any disclaimer, the term of this patent is extended or adjusted under 35 U.S.C. 154(b) by 125 days.

TW 1469047 B 1/2015

OTHER PUBLICATIONS

(21) Appl. No.: **13/993,508**

PCT International Search Report for PCT Counterpart Application No. PCT/US2011/068208, 3 pgs., (Sep. 27, 2012).

(22) PCT Filed: **Dec. 30, 2011**

(Continued)

(86) PCT No.: **PCT/US2011/068208**

§ 371 (c)(1),
(2), (4) Date: **Jun. 12, 2013**

Primary Examiner — Jared Rutz

Assistant Examiner — Jean Edouard

(74) *Attorney, Agent, or Firm* — Nicholson De Vos Webster & Elliott LLP

(87) PCT Pub. No.: **WO2013/101213**

PCT Pub. Date: **Jul. 4, 2013**

(57) **ABSTRACT**

In accordance with embodiments disclosed herein, there are provided methods, systems, mechanisms, techniques, and apparatuses for cutting senior store latency using store prefetching. For example, in one embodiment, such means may include an integrated circuit or an out of order processor means that processes out of order instructions and enforces in-order requirements for a cache. Such an integrated circuit or out of order processor means further includes means for receiving a store instruction; means for performing address generation and translation for the store instruction to calculate a physical address of the memory to be accessed by the store instruction; and means for executing a pre-fetch for a cache line based on the store instruction and the calculated physical address before the store instruction retires.

(65) **Prior Publication Data**

US 2014/0223105 A1 Aug. 7, 2014

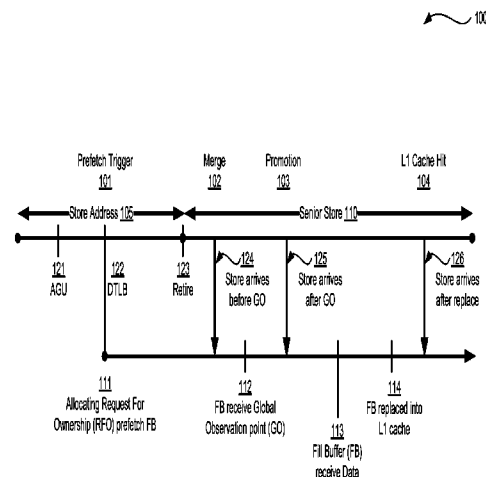
(51) **Int. Cl.**
G06F 12/08 (2016.01)
G06F 9/38 (2006.01)

(52) **U.S. Cl.**
CPC **G06F 9/3814** (2013.01); **G06F 9/383** (2013.01); **G06F 9/3834** (2013.01);
(Continued)

(58) **Field of Classification Search**

CPC G06F 12/0808; G06F 12/0862; G06F

22 Claims, 14 Drawing Sheets



(52) U.S. Cl.

CPC *G06F 9/3861* (2013.01); *G06F 12/0808*
(2013.01); *G06F 12/0862* (2013.01); *G06F*
2212/6028 (2013.01); *G06F 2212/62* (2013.01)

(56)

References Cited

U.S. PATENT DOCUMENTS

6,094,719 A * 7/2000 Panwar G06F 9/30014
708/513
6,216,215 B1 * 4/2001 Palanca et al. 712/23
6,266,744 B1 * 7/2001 Hughes G06F 9/3834
711/123
6,449,693 B1 * 9/2002 Goetz et al. 711/122
6,594,734 B1 * 7/2003 Kyker G06F 9/3812
711/125
6,643,745 B1 * 11/2003 Palanca G06F 9/383
711/137
6,742,082 B1 * 5/2004 Lango et al. 711/118
7,383,409 B2 6/2008 Steely, Jr. et al.
8,826,141 B2 * 9/2014 Dimitrov 715/738
8,832,288 B1 * 9/2014 Mahkovec et al. 709/229
2002/0035677 A1 * 3/2002 Sheaffer G06F 9/3836
712/217
2002/0056027 A1 * 5/2002 Kanai et al. 711/137
2002/0138714 A1 * 9/2002 Leibholz G06F 9/3836
712/217
2003/0084259 A1 * 5/2003 Palanca et al. 711/163

2003/0149862 A1 * 8/2003 Kadambi G06F 9/3838
712/217
2003/0149865 A1 * 8/2003 Kadambi G06F 9/3867
712/244
2004/0128448 A1 * 7/2004 Stark et al. 711/137
2006/0248319 A1 * 11/2006 Kadambi G06F 9/322
712/234
2008/0263331 A1 * 10/2008 Le et al. 712/217
2009/0157943 A1 * 6/2009 Robinson G06F 9/3834
711/100
2010/0106907 A1 * 4/2010 Noguchi et al. 711/114
2010/0146247 A1 * 6/2010 Ban G06F 9/3855
712/226
2010/0205415 A1 * 8/2010 Henry G06F 9/3867
712/234
2010/0306509 A1 * 12/2010 Day G06F 9/30174
712/217

OTHER PUBLICATIONS

PCT Written Opinion of the International Searching Authority for
PCT Counterpart Application No. PCT/US2011/068208, 5 pgs.,
(Sep. 27, 2012).

PCT Notification concerning Transmittal of International Prelimi-
nary Report on Patentability (Chapter I of the Patent Cooperation
Treaty) for PCT Counterpart Application No. PCT/US2011/068208,
7 pgs., (Jul. 10, 2014).

Taiwan IPO Search Report from counterpart application No. Taiwan
Patent Application No. 101148745, mailed Aug. 26, 2014.

* cited by examiner

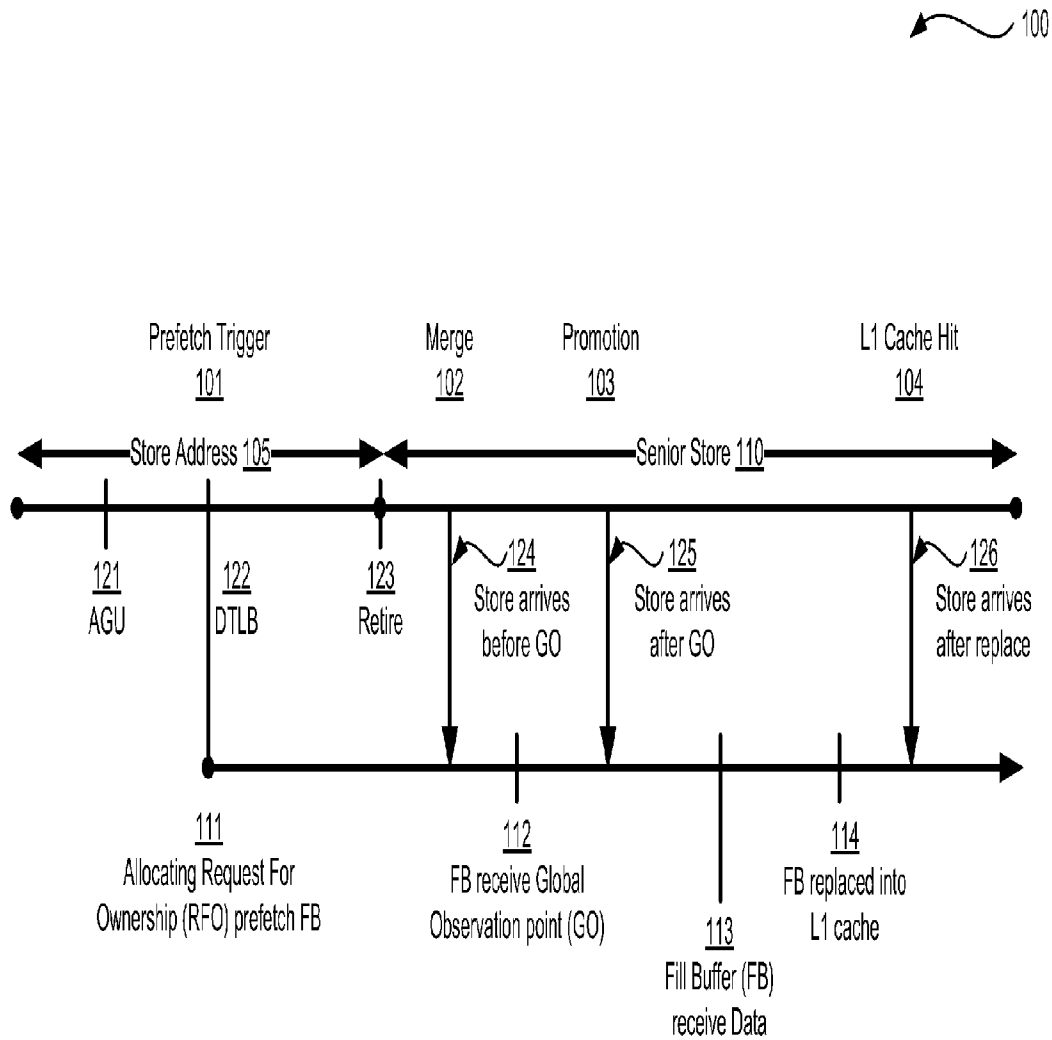


FIG. 1




FIG. 2




FIG. 3

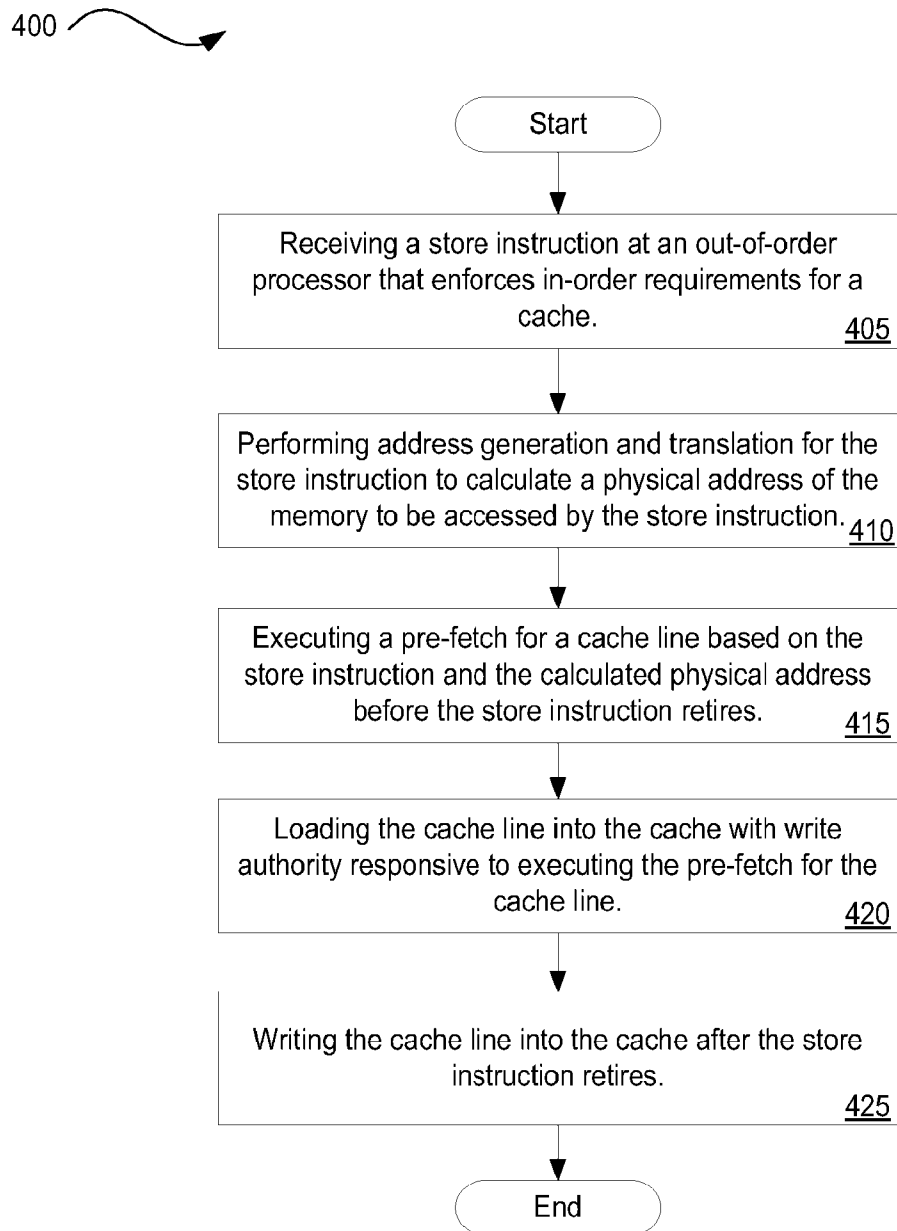


FIG. 4



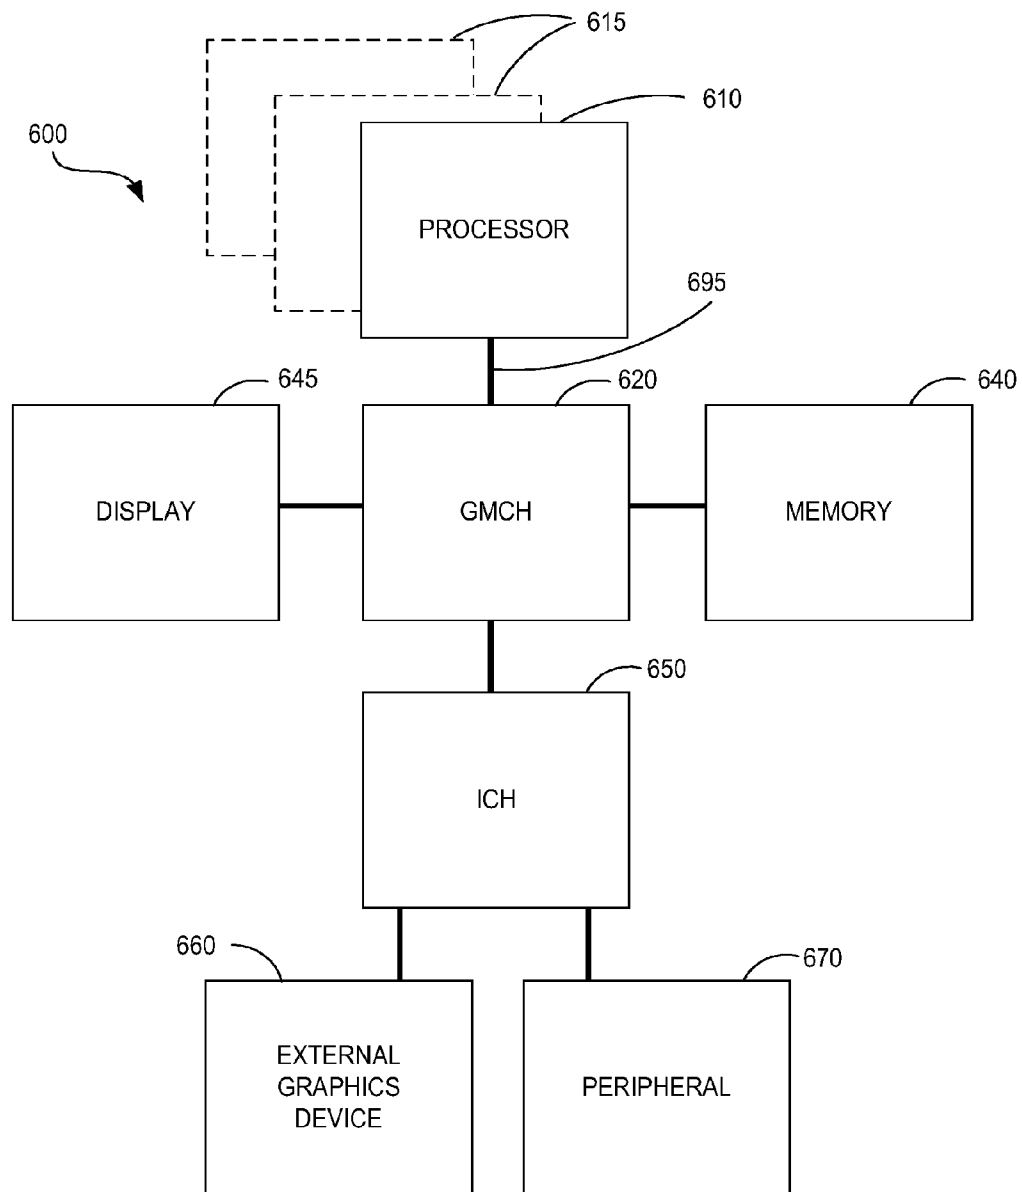
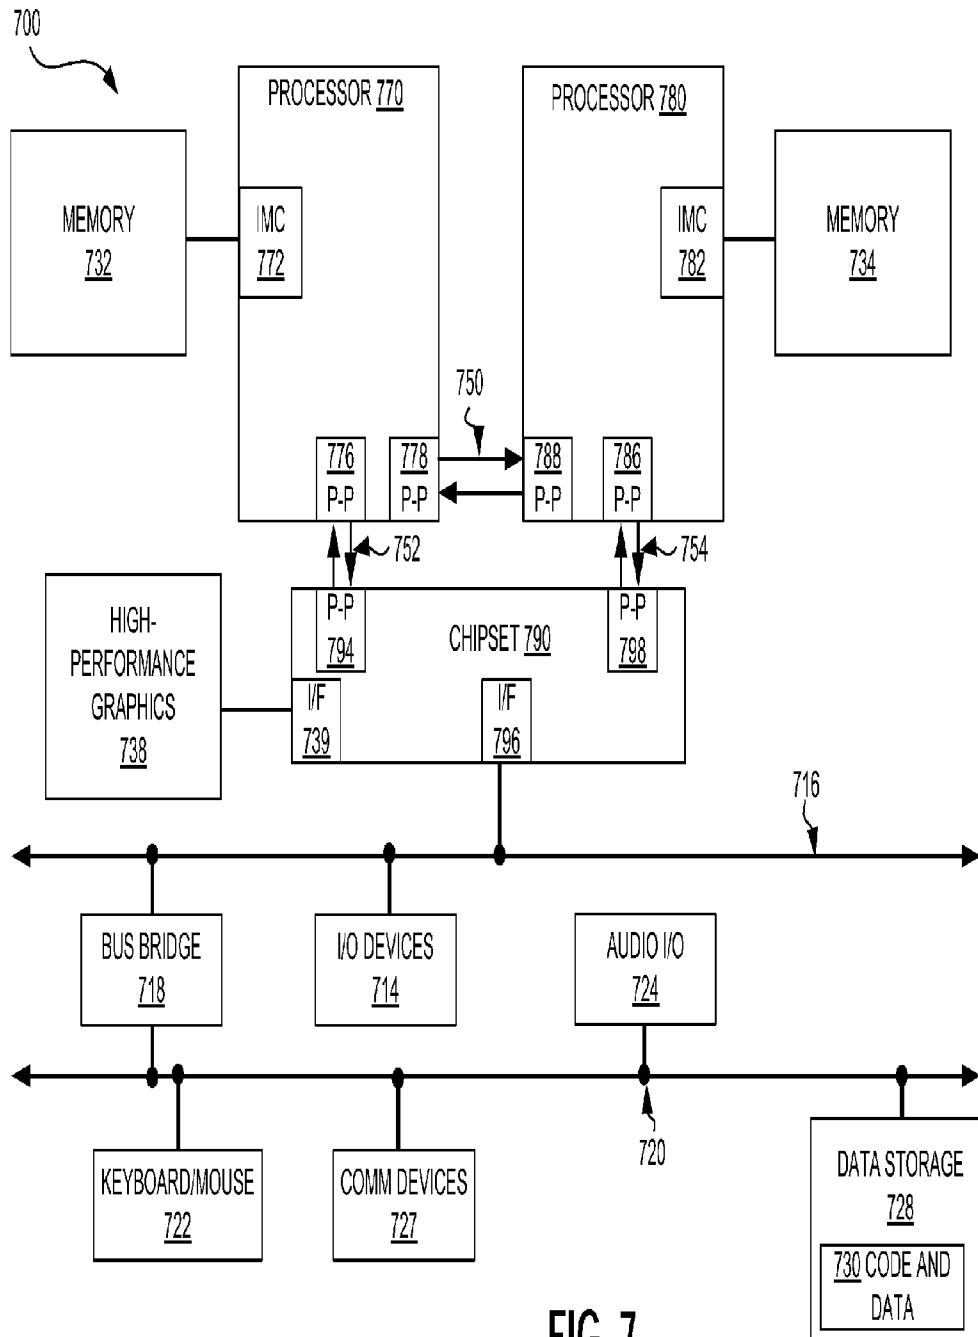


FIG. 5A

FIG. 5B



**FIG. 6**



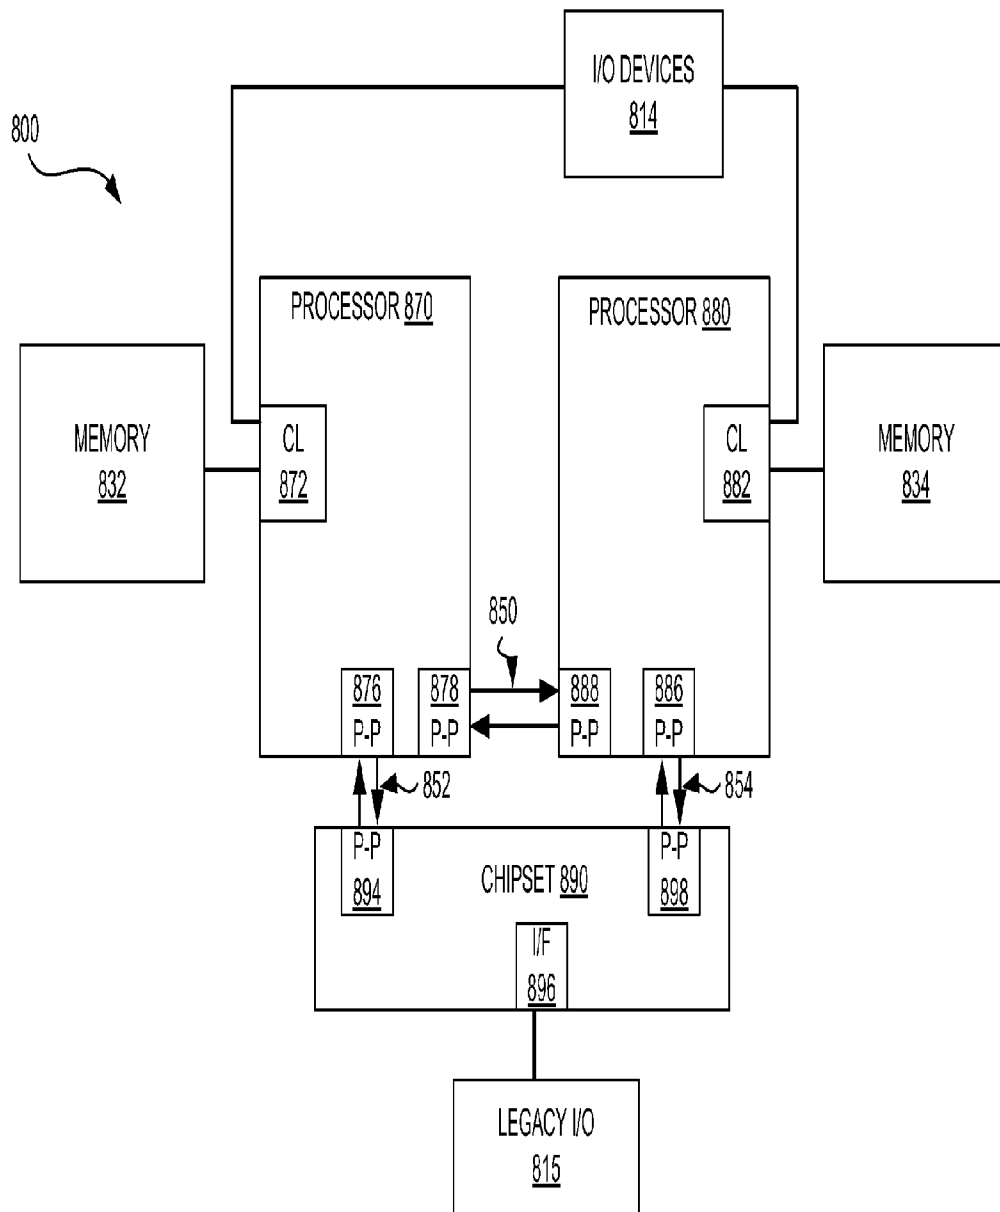


FIG. 8

FIG. 9


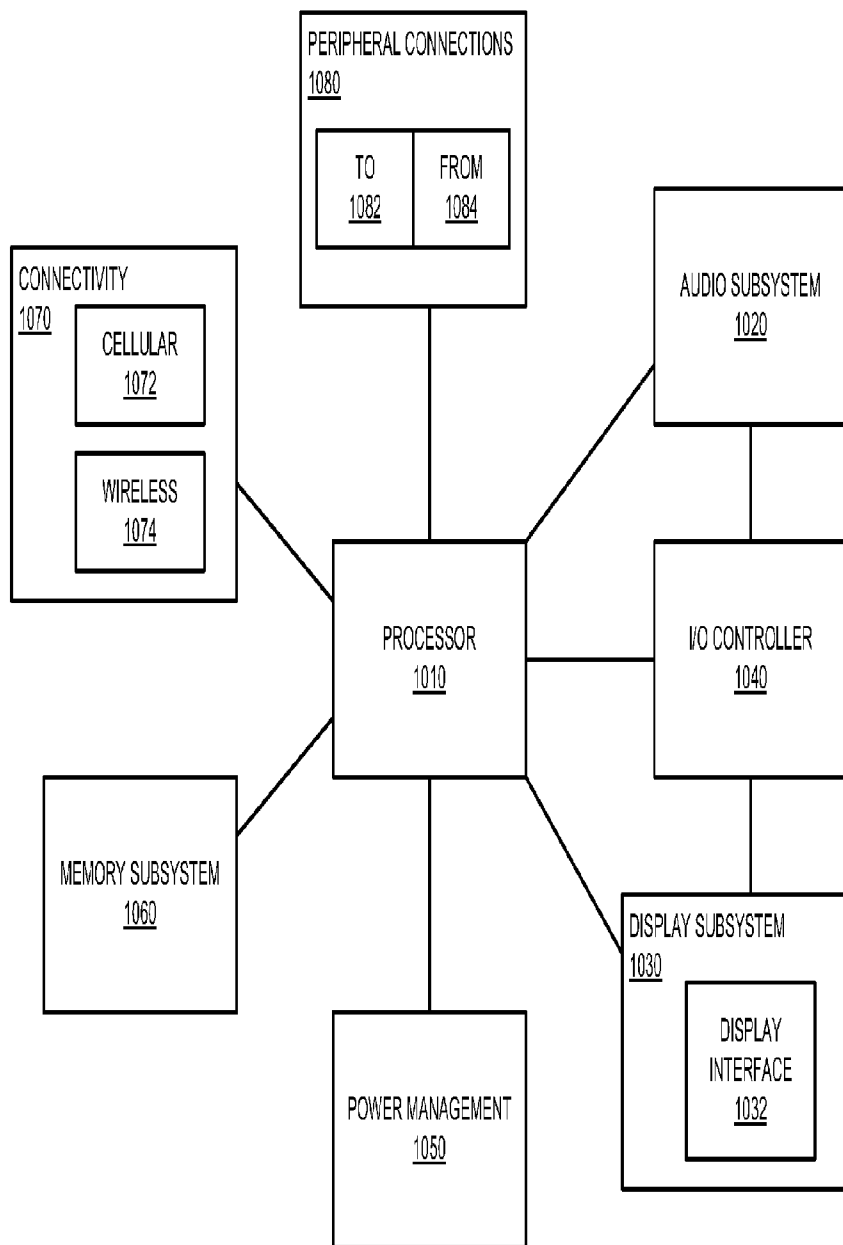


FIG. 10

1000





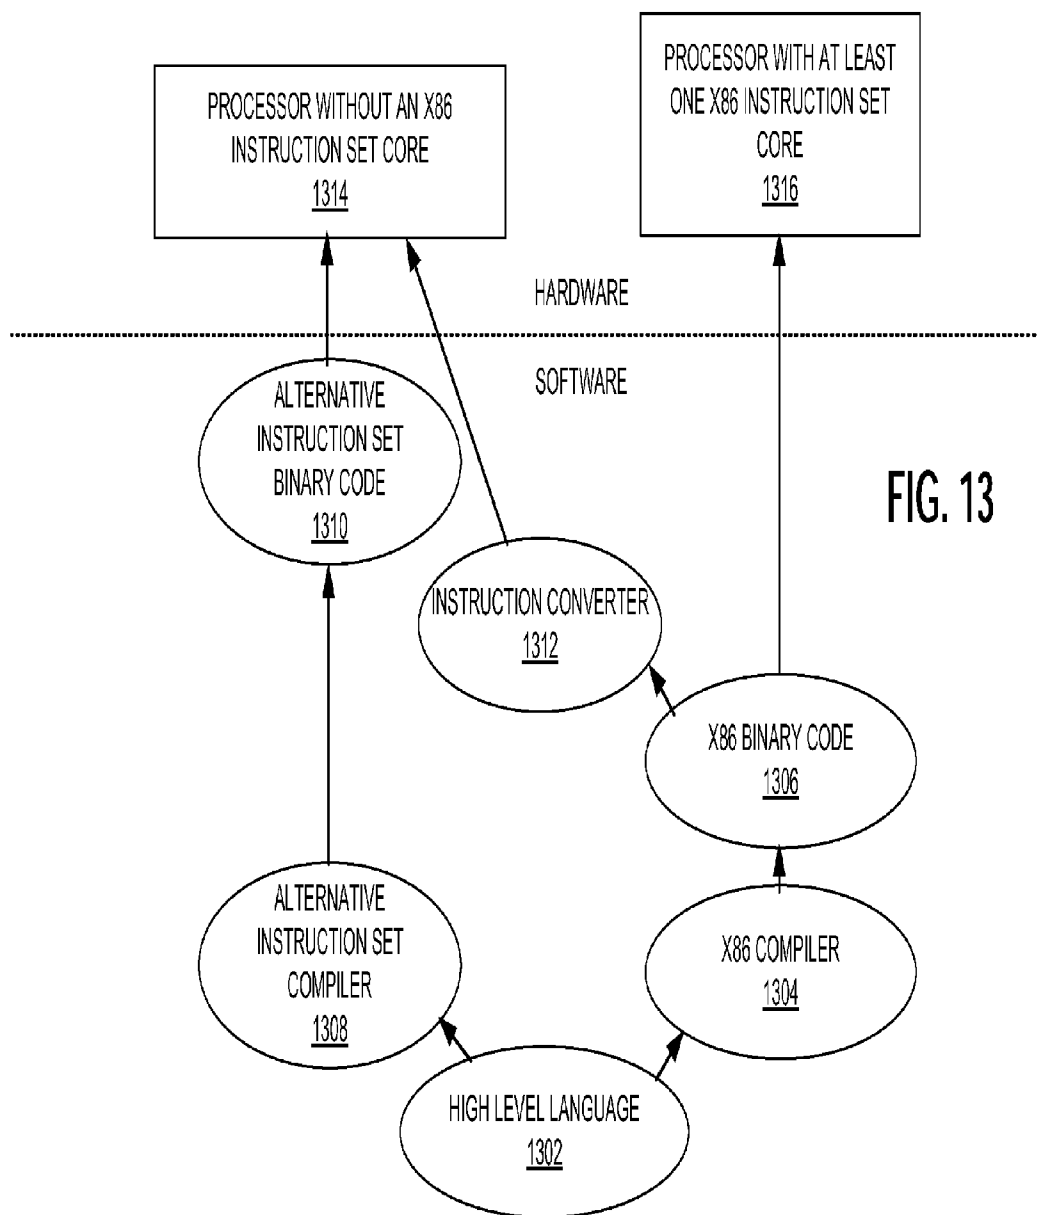


FIG. 11

**FIG. 12**



1

METHOD AND APPARATUS FOR CUTTING SENIOR STORE LATENCY USING STORE PREFETCHING

CROSS-REFERENCE TO RELATED APPLICATION

This patent application is a U.S. National Phase Application under 35 U.S.C. §371 of International Application No. PCT/US2011/068208, filed Dec. 30, 2011, entitled METHOD AND APPARATUS FOR CUTTING SENIOR STORE LATENCY USING STORE PREFETCHING.

COPYRIGHT NOTICE

A portion of the disclosure of this patent document contains material which is subject to copyright protection. The copyright owner has no objection to the facsimile reproduction by anyone of the patent document or the patent disclosure, as it appears in the Patent and Trademark Office patent file or records, but otherwise reserves all copyright rights whatsoever.

TECHNICAL FIELD

The subject matter described herein relates generally to the field of computing, and more particularly, to systems and methods for cutting senior store latency using store prefetching.

BACKGROUND

The subject matter discussed in the background section should not be assumed to be prior art merely as a result of its mention in the background section. Similarly, a problem mentioned in the background section or associated with the subject matter of the background section should not be assumed to have been previously recognized in the prior art. The subject matter in the background section merely represents different approaches, which in and of themselves may also correspond to embodiments of the claimed subject matter.

Generally speaking, memory closer to a CPU may be accessed faster than memory farther away. Memory within a CPU may be referred to as cache, and may be accessible at different hierarchical levels, such as Level 1 cache (L1 cache) and Level 2 cache (L2 cache). System memory such as memory modules coupled with a motherboard may also be available, such externally available memory which is separate from the CPU but accessible to the CPU may be referred to as, for example, off-chip cache or Level 3 cache (L3 cache), and so on, however, this is not always consistent as a third hierarchical level of cache (e.g., L3 cache) may be on-chip or “on-die” and thus be internal to the CPU.

CPU cache, such as L1 cache, is used by the central processing unit of a computer to reduce the average time to access memory. The L1 cache is a smaller, faster memory which stores copies of the data from the most frequently used main memory locations. L2 cache may be larger, but slower to access. And additional memory, whether on-chip or externally available system memory used as cache may be larger still, but slower to access than smaller and closer CPU cache levels. As long as most memory accesses are cached memory locations, the average latency of memory accesses will be closer to the cache latency than to the latency of main memory.

2

Conventional caches utilize a store buffer to reduce cache latency and also to enable the reading of store instructions that have not yet been written into cache. As stores go down a pipeline they store the data in a store buffer and persist until the store is retired from the pipeline, at which point the store writes the data to cache.

Conventional cache mechanisms require that store instructions be instituted through a series of operations which are executed in serial steps. Instructions are decoded and forwarded to an address generation unit, an address is calculated, and then sent to the cache which must maintain the order the instructions serially to carry out the store.

Moreover, the stored data cannot be made available to other entities until absolutely certain that the store is actually going to persist until retirement, at which point the store instruction “retires” from the pipeline thus allowing the stored data to be written from a store buffer to the cache location and it is at this post retirement stage of operation that the data is considered valid.

Because an out of order machine or processor executes instructions “out of order,” it cannot be known with certainty whether any given instruction will be part of a validly executed path. For example, where an instruction is executed ahead of a branch, there is a risk that such an instruction will never be used, should a branch mis-prediction occur. Thus, problems arise with data integrity if a store writes data into a cache before it is known with certainty whether or not the store instruction will retire from the cache, and thus correspond to valid data. A store instruction which never retires, yet writes to the cache, causes invalid data to be written to the cache and thus creates a data integrity problem. Unfortunately, retirement occurs at a late stage, and thus, induces cache latency for such store instructions.

Improvements to cache latency (e.g., reductions in cache latency) provide direct and immediate benefits to computational efficiency for an integrated circuit utilizing such a cache. Lower latency means that data required by, for example, a CPU pipeline is available sooner without having to expend cycles waiting for unavailable data.

The present state of the art may therefore benefit from systems and methods for cutting senior store latency using store prefetching as described herein.

BRIEF DESCRIPTION OF THE DRAWINGS

Embodiments are illustrated by way of example, and not by way of limitation, and will be more fully understood with reference to the following detailed description when considered in connection with the figures in which:

FIG. 1 illustrates an exemplary architecture in accordance with which embodiments may operate;

FIG. 2 illustrates an alternative exemplary architecture in accordance with which embodiments may operate;

FIG. 3 shows a diagrammatic representation of a system in accordance with which embodiments may operate, be installed, integrated, or configured;

FIG. 4 is a flow diagram illustrating a method for cutting senior store latency using store prefetching in accordance with described embodiments;

FIG. 5A is a block diagram illustrating both an exemplary in-order pipeline and an exemplary register renaming, out-of-order issue/execution pipeline in accordance with described embodiments;

FIG. 5B is a block diagram illustrating both an exemplary embodiment of an in-order architecture core and an exem-

3

plary register renaming, out-of-order issue/execution architecture core to be included in a processor in accordance with described embodiments;

FIG. 6 is a block diagram of a computer system according to one embodiment;

FIG. 7 is a block diagram of a computer system according to one embodiment;

FIG. 8 is a block diagram of a computer system according to one embodiment;

FIG. 9 depicts a tablet computing device and a hand-held smartphone each having a circuitry integrated therein as described in accordance with the embodiments;

FIG. 10 is a block diagram of an embodiment of tablet computing device, a smartphone, or other mobile device in which touchscreen interface connectors are used;

FIG. 11 is a block diagram of an IP core development system according to one embodiment;

FIG. 12 illustrates an architecture emulation system according to one embodiment; and

FIG. 13 illustrates a system to translate instructions according to one embodiment.

DETAILED DESCRIPTION

Described herein are systems and methods for cutting senior store latency using store prefetching. For example, in one embodiment, such means may include an integrated circuit or an out of order processor means that processes out-of-order instructions and enforces in-order requirements for a cache. Such an integrated circuit or out of order processor means further includes means for receiving a store instruction; means for performing address generation and translation for the store instruction to calculate a physical address of the memory to be accessed by the store instruction; and means for executing a pre-fetch for a cache line based on the store instruction and the calculated physical address before the store instruction retires.

In the following description, numerous specific details are set forth such as examples of specific systems, languages, components, etc., in order to provide a thorough understanding of the various embodiments. It will be apparent, however, to one skilled in the art that these specific details need not be employed to practice the embodiments disclosed herein. In other instances, well known materials or methods have not been described in detail in order to avoid unnecessarily obscuring the disclosed embodiments.

In addition to various hardware components depicted in the figures and described herein, embodiments further include various operations which are described below. The operations described in accordance with such embodiments may be performed by hardware components or may be embodied in machine-executable instructions, which may be used to cause a general-purpose or special-purpose processor programmed with the instructions to perform the operations. Alternatively, the operations may be performed by a combination of hardware and software.

Embodiments also relate to an apparatus for performing the operations disclosed herein. This apparatus may be specially constructed for the required purposes, or it may be a general purpose computer selectively activated or reconfigured by a computer program stored in the computer. Such a computer program may be stored in a computer readable storage medium, such as, but not limited to, any type of disk including floppy disks, optical disks, CD-ROMs, and magnetic-optical disks, read-only memories (ROMs), random access memories (RAMs), EPROMs, EEPROMs, magnetic or optical cards, or any type of media suitable for storing

4

electronic instructions, each coupled with a computer system bus. The term “coupled” may refer to two or more elements which are in direct contact (physically, electrically, magnetically, optically, etc.) or to two or more elements that are not in direct contact with each other, but still cooperate and/or interact with each other.

The algorithms and displays presented herein are not inherently related to any particular computer or other apparatus. Various general purpose systems may be used with programs in accordance with the teachings herein, or it may prove convenient to construct more specialized apparatus to perform the required method steps. The required structure for a variety of these systems will appear as set forth in the description below. In addition, embodiments are not described with reference to any particular programming language. It will be appreciated that a variety of programming languages may be used to implement the teachings of the embodiments as described herein.

Any of the disclosed embodiments may be used alone or together with one another in any combination. Although various embodiments may have been partially motivated by deficiencies with conventional techniques and approaches, some of which are described or alluded to within the specification, the embodiments need not necessarily address or solve any of these deficiencies, but rather, may address only some of the deficiencies, address none of the deficiencies, or be directed toward different deficiencies and problems which are not directly discussed.

FIG. 1 illustrates an exemplary architecture **100** in accordance with which embodiments may operate.

In modern out-of-order micro-architectures store instructions must nevertheless be executed in-order so as to maintain data integrity which could arise should a store instruction cause data to be written to cache out of order and then fail to retire, thus rendering the data written to cache as erroneous. Such an issue occurs because 1) a write operation to the cache is destructive and cannot be rolled back in the case of a mis-speculation; and 2) conventional x86 architecture memory ordering rules require that store operations be visible strictly in-order.

Actual store writes to the cache are not permitted until the store instruction is retired and the store becomes senior. Described mechanisms to reduce senior store lifetime are therefore proposed which in turn improves cache latency and overall performance. A senior store **110** means that the store instruction has passed retirement and is no longer speculative as it is known with certainty that the store instruction is going to execute (e.g., software has already executed the instruction) and thus, the data must be written to cache to ensure data integrity.

Senior store lifetime is reduced by prefetching ahead the data that a senior store is going to access. Rather than issuing a Request For Ownership (RFO) request after the arrival of a senior store to the cache, an RFO prefetch request is issued during the store address execution time immediately after the store instruction physical address is known. For example, with reference to the exemplary architecture **100** of FIG. 1, a “store address” **105** phase is followed by the “senior store” **110** phase. The RFO prefetch request may be issued using a hardware prefetching engine or issued directly to the cache.

Some cache mechanisms may lack sufficient distance between the RFO prefetch and the senior store **110** such that the senior store would reach the Fill Buffer (FB) processing the RFO prefetch request before the RFO prefetch completes. Therefore, in accordance with the described embodiments, a promotion **103** mechanism is further disclosed which converts a Fill Buffer (FB) allocated for the RFO prefetch into a

5

demand senior store FB while maintaining strong x86 memory ordering requirements, thus allowing the senior store **110** to complete earlier and deallocate its store buffer entry. The disclosed technique does not attempt to guess in advance what the next address a store is going to access (as some variants of a pre-fetcher may), and thus, there is no risk of a mis-speculation.

Various events and their relative relationships to each other are depicted by the exemplary architecture **100** of FIG. 1. Within the store address **105** phase, Address Generation Unit (AGU) **121** precedes the pre-trigger **101** event which corresponds with the allocating RFO prefetch FB **111**. Data Translation Lookaside Buffer (DTLB) **122** event and Retire **123** complete the store address **105** phase. Merge **102** is depicted in the senior store **110** phase as is promotion **103** followed by an L1 cache hit **104**. Event **124** corresponding to a store arrives before Global Observation point (GO) event precedes a FB receives GO event **112**, followed by a store arrives after GO event **125** which precedes the FB receive data event **113** and the FB replaced into L1 cache event **114**, which is followed by the store arrives after replace event **126**.

If the store instruction writes to cache immediately after retirement, then operation is already after the Global Observation point, and thus, the store instruction may write to the cache without special handling. The allocating the RFO prefetch FB event **111** performs a prefetch store request based on a physical address to get write authority, and thus establishes an independent pipeline for the prefetch. After a time, the senior store **110** arrives. If the senior store **110** arrives before the prefetch request Global Observation point then merge event **102** is initiated. If the senior store **110** arrives after prefetch has already received its Global Observation point then promotion event **103** is initiated. If the senior store **110** arrives after prefetch is replaced, then no special handling is required.

FIG. 2 illustrates another alternative exemplary architecture **200** in accordance with which embodiments may operate. For example, an integrated circuit **201** is depicted as having pipeline **299**, store buffer **205**, cache **210**, fill buffer **235** with attribute **236**, Address Generation Unit (AGU) **225**, each of which are communicably interfaced to the data bus **215** of integrated circuit **201**.

In accordance with one embodiment, the integrated circuit **201** implements an out-of-order processor. In such an embodiment, a pipeline **299** the out-of-order processor receives a store instruction. In such an embodiment, the out-of-order processor enforces in-order requirements for a cache **210**. The out-of-order processor further performing address generation and translation for the store instruction to calculate a physical address of the memory to be accessed by the store instruction, for example, via an Address Generation Unit **225** and executing a pre-fetch for a cache line based on the store instruction and the calculated physical address before the store instruction retires, for example, via fill buffer **235**.

In one embodiment, the integrated circuit **201** includes a Central Processing Unit (CPU). In one embodiment, the central processing unit operates as a CPU or an out-of-order processor CPU for one of a tablet computing device or a smartphone.

In one embodiment, receiving the store instruction at the out-of-order processor includes receiving the store instruction at the pipeline **299** of the out-of-order processor.

In one embodiment, the integrated circuit **201** executes the pre-fetch for the cache line before the store instruction retires by executing the pre-fetch for the cache line before the store instruction retires from the pipeline **299**. In one embodiment,

6

the store instruction is subject to one of the following events: (a) retirement from the pipeline **299** pursuant to an accurate instruction prediction or (b) a pipeline flush event pursuant to an instruction mis-prediction. In one embodiment, the pipeline flush event results in the store instruction being rendered obsolete. For example, if the store instruction is flushed from the pipeline before retirement, it will never retire, and thus, the data corresponding to the store instruction will not be valid.

In one embodiment, the fill buffer executes the pre-fetch for the cache line by requesting write authority to the cache line within the cache **210**. In one embodiment, the cache **210** includes a Level 1 cache (L1 cache); and executing the pre-fetch for the cache line results in a cache miss at the L1 cache. In such an embodiment, executing the pre-fetch for the cache line further includes one or more of: (a) executing the pre-fetch against a higher level cache (e.g., such as an Level 2 cache, a Level 3 cache, etc.) responsive to the cache miss; or (b) executing the pre-fetch against a system non-cache random access memory (RAM) communicably interfaced with the out-of-order processor responsive to the cache miss. For example, there may be no cache level of integrated circuit **201** that contains the cache line, and thus, the data must be retrieved from system memory or RAM, and in particular, non-cache RAM. In one embodiment, the fill buffer **235** further loads the cache line into the cache **210** with write authority responsive to executing the pre-fetch for the cache line.

In order to reduce latency on a cache miss, it is advantageous to initiate a pre-fetch trigger (as shown at event **101** of FIG. 1) for the required cache line. Issuing the pre-fetch operation (e.g., allocating RFO prefetch FB at event **111**) in advance of the store instruction retiring from the pipeline allows the required cache line to be retrieved and made available in the cache **210** before the store instruction would ordinarily even attempt to perform the cache write. Such a technique can be especially advantageous if the pre-fetch must go all the way out to system RAM at which point latency will be about 200 cycles. The pre-fetch leverages a trick of timing to cause the retrieval of the cache line which in turn reduces latency for the eventual cache write to the cache **210** upon retirement of the store instruction to approximately 5 cycles. Without the pre-retirement pre-fetch of the cache line, a cache miss would require the cache write to wait (inducing latency) until the cache line is retrieved in a serial fashion subsequent to retirement of the store instruction, then execute an RFO to attain write authority, and finally conduct the actual write to the cache **210**.

Most pre-fetch algorithms attempt to perform a predictive guess of what data will be needed and then issue a reads for the corresponding addresses based on the predictive guess. Unfortunately, such a mechanism is not always accurate and is therefore subject to the severe latency associated with a cache miss and then a serial retrieval, RFO request, and eventual cache write.

In accordance with one embodiment, the store buffer **205** writes the cache line into the cache **210** after the store instruction retires and deallocating the store buffer **205** allocated to, or associated with, the cache line. Because the cache line is written to cache **210** only after retirement, the actual address is known well in advance of attempting the cache write. The store address is known in the pipeline **299** very early, specifically, at execution time (refer to store address **105** phase of FIG. 1). Execution always occurs in advance of retirement and sometimes occurs much earlier, tens to potentially hundreds of cycles earlier. Accordingly, the physical address may be utilized, without having to rely upon predictive guessing, to execute the pre-fetch operation.

In certain situations, pre-fetch alone is not sufficient due to a necessity to enforce ordering requirements for the integrated circuit **201**. In an effort to complete the cache write and deallocate from the store buffer **205** as quickly as possible, the pre-fetch is issued as early as feasible, such as immediately after the store instruction physical address is known. However, if the pre-fetch is not yet complete when the senior store (e.g., see senior store **110** phase of FIG. 1) arrives, and the pre-fetch is still working to bring the required cache line into the cache **210**, then the store buffer **205** cannot be deallocated.

Thus, in accordance with one embodiment, the integrated circuit **201** further (a) identifies a pre-existing pre-fetch operation for the same cache line and (b) promotes the pre-existing pre-fetch operation to a senior store. In such an embodiment, the integrated circuit **201** promotes the pre-existing pre-fetch operation to the senior store by (a) identifying a fill buffer **235** allocated to bring the cache line from a memory (e.g., a higher cache or separate system memory) into the cache **210** pursuant to the pre-existing pre-fetch operation; and (b) by further modifying an attribute **236** of the fill buffer **235** from indicating a pre-fetch operation to indicating a senior store operation. In such an embodiment, promoting the pre-existing pre-fetch operation to the senior store fulfills the pre-fetch for the cache line.

The promotion mechanism (e.g., see promotion event **103** of FIG. 1) provides a special handling mechanism to evaluate a store that comes to the cache **210** and encounters an already existing prefetch to bring to the same line, thus resulting in a contention problem. So as to overcome the contention, the fill buffer **235** working on retrieving the same cache line is taken over and converted from the pre-existing pre-fetch operation and into a senior store by modifying the appropriate attribute **236** of the fill buffer. Modification of attribute **236** thus causes the buffer to appear and operate as though it had been a senior store operation from the outset.

Thus, in accordance with a particular embodiment, the integrated circuit **201** promotes the pre-existing pre-fetch operation to the senior store by (a) determining the pre-fetch for the cache line has not completed; (b) by further identifying a fill buffer **235** allocated to bring the cache line from a memory into the cache **210** pursuant to the pre-existing pre-fetch operation; and (c) by further overwriting an attribute of the fill buffer **235** from indicating the pre-existing pre-fetch operation with an attribute indicating a senior store. In such an embodiment, promoting the pre-existing pre-fetch operation to the senior store again fulfills the pre-fetch for the cache line.

In accordance with one embodiment, the out-of-order processor or integrated circuit **201** enforces the in-order requirements for a cache **210** by requiring that stores to the cache **210** in fulfillment of store instructions received at the out-of-order processor are made visible external to the out-of-order processor in strict order. In such an embodiment, the strict order is identical in sequence to a program order being executed by the out-of-order processor. That is, the order of execution of, for example, an application or other software will correspond to the order in which store instructions retire from the pipeline **299** and in which stores are eventually written to cache **210**. In one embodiment, a real linked list of the order in the fill buffer **235** is maintained in support of enforcing the in-order cache requirements.

In accordance with one embodiment, the out-of-order processor or integrated circuit **201** further (a) determines the pre-fetch for the cache line has not completed; (b) determines a Global Observation point (GO) has been reached for the pre-fetch; and (c) hides a Global Observation point from being visible external to the out-of-order processor prevent-

ing access to the cache line from entities external to the out-of-order processor. In such an embodiment, the out-of-order processor or integrated circuit **201** may further (d) determine all previous senior stores are associated with a corresponding Global Observation point and (e) un-hide (e.g., make visible) the Global Observation point for the pre-fetch allowing the pre-fetch to receive a Global Observation point time-stamp subsequent to all previous stores to ensure in-order requirements of the out-of-order processor.

Thus, a GO point is reached but the pre-fetch has not yet completed, then guarantees must be instituted to ensure compliance with strict order with every previous store. Thus, GO point is hidden from the main memory, from the program, etc., until such time that all previous senior stores receive their own GO. After such time, the pre-fetch operation will be allowed to receive its own GO time stamp and become visible. Thus, even if a GO is feasibly available, the GO will not be strict order compliant, and thus, it is hidden from entities external to the out-of-order processor or integrated circuit **201** responsible for maintaining ordering compliance.

FIG. 3 shows a diagrammatic representation of a system **399** in accordance with which embodiments may operate, be installed, integrated, or configured.

In one embodiment, system **399** includes a memory **395** and a processor or processors **396**. For example, memory **395** may store instructions to be executed and processor(s) **396** may execute such instructions. System **399** includes communication bus(es) **365** to transfer transactions, instructions, requests, and data within system **399** among a plurality of peripheral device(s) **370** communicably interfaced with one or more communication buses **365** and/or interface(s) **375**. Display unit **380** is additionally depicted within system **399**.

Distinct within system **399** is integrated circuit **301** which may be installed and configured in a compatible system **399**, or manufactured and provided separately so as to operate in conjunction with appropriate components of system **399**.

In accordance with one embodiment, system **399** includes at least a display unit **380** and an integrated circuit **301**. The integrated circuit **301** may operate as, for example, an out-of-order processor or as another computing component of system **399**. In such an embodiment, the integrated circuit **301** of system **399** includes at least: a data bus **315**, a pipeline **305** to receive a store instruction in which the integrated circuit is to process out-of-order instructions and enforce in-order requirements for a cache **310**; an Address Generation Unit (AGU) **325** to calculate a physical address for the store instruction; and a fill buffer **335** to pre-fetch a cache line based on the store instruction and the calculated physical address before the store instruction retires from the pipeline **305**.

In accordance with one embodiment, such a system **399** embodies a tablet or a smartphone, in which the display unit **380** is a touchscreen interface of the tablet or the smartphone; and further in which the integrated circuit **301** is incorporated into the tablet or smartphone.

FIG. 4 is a flow diagram illustrating a method for cutting senior store latency using store prefetching in accordance with described embodiments. Method **400** may be performed by processing logic that may include hardware (e.g., circuitry, dedicated logic, programmable logic, microcode, etc.). The numbering of the blocks presented is for the sake of clarity and is not intended to prescribe an order of operations in which the various blocks must occur.

Method **400** begins with processing logic for receiving a store instruction at an out-of-order processor that enforces in-order requirements for a cache (block **405**).

At block **410**, processing logic performs address generation and translation for the store instruction to calculate a physical address of the memory to be accessed by the store instruction.

At block **415**, processing logic executes a pre-fetch for a cache line based on the store instruction and the calculated physical address before the store instruction retires.

At block **420**, processing logic loads the cache line into the cache with write authority responsive to executing the pre-fetch for the cache line.

And at block **425**, processing logic writes the cache line into the cache after the store instruction retires.

FIG. **5A** is a block diagram illustrating both an exemplary in-order pipeline and an exemplary register renaming, out-of-order issue/execution pipeline in accordance with described embodiments. FIG. **5B** is a block diagram illustrating both an exemplary embodiment of an in-order architecture core and an exemplary register renaming, out-of-order issue/execution architecture core to be included in a processor in accordance with described embodiments. The solid lined boxes in FIG. **5A** and FIG. **5B** illustrate the in-order pipeline and in-order core, while the optional addition of the dashed lined boxes illustrates the register renaming, out-of-order issue/execution pipeline and core. Given that the in-order aspect is a subset of the out-of-order aspect, the out-of-order aspect will be described.

Processor cores may be implemented in different ways, for different purposes, and in different processors. For instance, implementations of such cores may include: 1) a general purpose in-order core intended for general-purpose computing; 2) a high performance general purpose out-of-order core intended for general-purpose computing; 3) a special purpose core intended primarily for graphics and/or scientific (throughput) computing. Implementations of different processors may include: 1) a CPU including one or more general purpose in-order cores intended for general-purpose computing and/or one or more general purpose out-of-order cores intended for general-purpose computing; and 2) a coprocessor including one or more special purpose cores intended primarily for graphics and/or scientific (throughput). Such different processors lead to different computer system architectures, which may include: 1) the coprocessor on a separate chip from the CPU; 2) the coprocessor on a separate die in the same package as a CPU; 3) the coprocessor on the same die as a CPU (in which case, such a coprocessor is sometimes referred to as special purpose logic, such as integrated graphics and/or scientific (throughput) logic, or as special purpose cores); and 4) a system on a chip that may include on the same die the described CPU (sometimes referred to as the application core(s) or application processor(s)), the above described coprocessor, and additional functionality. Exemplary core architectures are described next, followed by descriptions of exemplary processors and computer architectures.

In FIG. **5A**, a processor pipeline **500** includes a fetch stage **502**, a length decode stage **504**, a decode stage **506**, an allocation stage **508**, a renaming stage **510**, a scheduling (also known as a dispatch or issue) stage **512**, a register read/memory read stage **514**, an execute stage **516**, a write back/memory write stage **518**, an exception handling stage **522**, and a commit stage **524**.

FIG. **5B** shows processor core **590** including a front end unit **530** coupled to an execution engine unit **550**, and both are coupled to a memory unit **570**. The core **590** may be a reduced instruction set computing (RISC) core, a complex instruction set computing (CISC) core, a very long instruction word (VLIW) core, or a hybrid or alternative core type. As yet another option, the core **590** may be a special-purpose core,

such as, for example, a network or communication core, compression engine, coprocessor core, general purpose computing graphics processing unit (GPGPU) core, graphics core, or the like.

The front end unit **530** includes a branch prediction unit **532** coupled to an instruction cache unit **534**, which is coupled to an instruction translation lookaside buffer (TLB) **536**, which is coupled to an instruction fetch unit **538**, which is coupled to a decode unit **540**. The decode unit **540** (or decoder) may decode instructions, and generate as an output one or more micro-operations, micro-code entry points, microinstructions, other instructions, or other control signals, which are decoded from, or which otherwise reflect, or are derived from, the original instructions. The decode unit **540** may be implemented using various different mechanisms. Examples of suitable mechanisms include, but are not limited to, look-up tables, hardware implementations, programmable logic arrays (PLAs), microcode read only memories (ROMs), etc. In one embodiment, the core **590** includes a microcode ROM or other medium that stores microcode for certain macroinstructions (e.g., in decode unit **540** or otherwise within the front end unit **530**). The decode unit **540** is coupled to a rename/allocator unit **552** in the execution engine unit **550**.

The execution engine unit **550** includes the rename/allocator unit **552** coupled to a retirement unit **554** and a set of one or more scheduler unit(s) **556**. The scheduler unit(s) **556** represents any number of different schedulers, including reservations stations, central instruction window, etc. The scheduler unit(s) **556** is coupled to the physical register file(s) unit(s) **558**. Each of the physical register file(s) units **558** represents one or more physical register files, different ones of which store one or more different data types, such as scalar integer, scalar floating point, packed integer, packed floating point, vector integer, vector floating point, status (e.g., an instruction pointer that is the address of the next instruction to be executed), etc. In one embodiment, the physical register file(s) unit **558** comprises a vector registers unit, a write mask registers unit, and a scalar registers unit. These register units may provide architectural vector registers, vector mask registers, and general purpose registers. The physical register file(s) unit(s) **558** is overlapped by the retirement unit **554** to illustrate various ways in which register renaming and out-of-order execution may be implemented (e.g., using a reorder buffer(s) and a retirement register file(s); using a future file(s), a history buffer(s), and a retirement register file(s); using a register maps and a pool of registers; etc.). The retirement unit **554** and the physical register file(s) unit(s) **558** are coupled to the execution cluster(s) **560**. The execution cluster(s) **560** includes a set of one or more execution units **562** and a set of one or more memory access units **564**. The execution units **562** may perform various operations (e.g., shifts, addition, subtraction, multiplication) and on various types of data (e.g., scalar floating point, packed integer, packed floating point, vector integer, vector floating point). While some embodiments may include a number of execution units dedicated to specific functions or sets of functions, other embodiments may include only one execution unit or multiple execution units that all perform all functions. The scheduler unit(s) **556**, physical register file(s) unit(s) **558**, and execution cluster(s) **560** are shown as being possibly plural because certain embodiments create separate pipelines for certain types of data/operations (e.g., a scalar integer pipeline, a scalar floating point/packed integer/packed floating point/vector integer/vector floating point pipeline, and/or a memory access pipeline that each have their own scheduler unit, physical register file(s) unit, and/or execution cluster—and in the case of a separate memory access pipeline, certain embodiments are

implemented in which only the execution cluster of this pipeline has the memory access unit(s) **564**). It should also be understood that where separate pipelines are used, one or more of these pipelines may be out-of-order issue/execution and the rest in-order.

The set of memory access units **564** is coupled to the memory unit **570**, which includes a data TLB unit **572** coupled to a data cache unit **574** coupled to a level 2 (L2) cache unit **576**. In one exemplary embodiment, the memory access units **564** may include a load unit, a store address unit, and a store data unit, each of which is coupled to the data TLB unit **572** in the memory unit **570**. The instruction cache unit **534** is further coupled to a level 2 (L2) cache unit **576** in the memory unit **570**. The L2 cache unit **576** is coupled to one or more other levels of cache and eventually to a main memory.

By way of example, the exemplary register renaming, out-of-order issue/execution core architecture may implement the pipeline **500** as follows: 1) the instruction fetch **538** performs the fetch and length decoding stages **502** and **504**; 2) the decode unit **540** performs the decode stage **506**; 3) the rename/allocator unit **552** performs the allocation stage **508** and renaming stage **510**; 4) the scheduler unit(s) **556** performs the schedule stage **512**; 5) the physical register file(s) unit(s) **558** and the memory unit **570** perform the register read/memory read stage **514**; the execution cluster **560** perform the execute stage **516**; 6) the memory unit **570** and the physical register file(s) unit(s) **558** perform the write back/memory write stage **518**; 7) various units may be involved in the exception handling stage **522**; and 8) the retirement unit **554** and the physical register file(s) unit(s) **558** perform the commit stage **524**.

The core **590** may support one or more instructions sets (e.g., the x86 instruction set (with some extensions that have been added with newer versions); the MIPS instruction set of MIPS Technologies of Sunnyvale, Calif.; the ARM instruction set (with optional additional extensions such as NEON) of ARM Holdings of Sunnyvale, Calif.), including the instruction(s) described herein. In one embodiment, the core **590** includes logic to support a packed data instruction set extension (e.g., AVX1, AVX2), thereby allowing the operations used by many multimedia applications to be performed using packed data.

It should be understood that the core may support multithreading (executing two or more parallel sets of operations or threads), and may do so in a variety of ways including time sliced multithreading, simultaneous multithreading (where a single physical core provides a logical core for each of the threads that physical core is simultaneously multithreading), or a combination thereof (e.g., time sliced fetching and decoding and simultaneous multithreading thereafter such as in the Intel® Hyperthreading technology).

While register renaming is described in the context of out-of-order execution, it should be understood that register renaming may be used by an in-order architecture. While the illustrated embodiment of the processor also includes separate instruction and data cache units **534/574** and a shared L2 cache unit **576**, alternative embodiments may have a single internal cache for both instructions and data, such as, for example, a Level 1 (L1) internal cache, or multiple levels of internal cache. In some embodiments, the system may include a combination of an internal cache and an external cache that is external to the core and/or the processor. Alternatively, all of the cache may be external to the core and/or the processor.

Referring now to FIG. 6, shown is a block diagram of a system **600** in accordance with one embodiment of the present invention. The system **600** may include one or more

processors **610**, **615**, which are coupled to graphics memory controller hub (GMCH) **620**. The optional nature of additional processors **615** is denoted in FIG. 6 with broken lines.

Each processor **610**, **615** may be some version of the circuit, integrated circuit, processor, and/or silicon integrated circuit as described above. However, it should be noted that it is unlikely that integrated graphics logic and integrated memory control units would exist in the processors **610**, **615**. FIG. 6 illustrates that the GMCH **620** may be coupled to a memory **640** that may be, for example, a dynamic random access memory (DRAM). The DRAM may, for at least one embodiment, be associated with a non-volatile cache.

The GMCH **620** may be a chipset, or a portion of a chipset. The GMCH **620** may communicate with the processor(s) **610**, **615** and control interaction between the processor(s) **610**, **615** and memory **640**. The GMCH **620** may also act as an accelerated bus interface between the processor(s) **610**, **615** and other elements of the system **600**. For at least one embodiment, the GMCH **620** communicates with the processor(s) **610**, **615** via a multi-drop bus, such as a frontside bus (FSB) **695**.

Furthermore, GMCH **620** is coupled to a display **645** (such as a flat panel or touchscreen display). GMCH **620** may include an integrated graphics accelerator. GMCH **620** is further coupled to an input/output (I/O) controller hub (ICH) **650**, which may be used to couple various peripheral devices to system **600**. Shown for example in the embodiment of FIG. 6 is an external graphics device **660**, which may be a discrete graphics device coupled to ICH **650**, along with another peripheral device **670**.

Alternatively, additional or different processors may also be present in the system **600**. For example, additional processor(s) **615** may include additional processors(s) that are the same as processor **610**, additional processor(s) that are heterogeneous or asymmetric to processor **610**, accelerators (such as, e.g., graphics accelerators or digital signal processing (DSP) units), field programmable gate arrays, or any other processor. There can be a variety of differences between the processor(s) **610**, **615** in terms of a spectrum of metrics of merit including architectural, micro-architectural, thermal, power consumption characteristics, and the like. These differences may effectively manifest themselves as asymmetry and heterogeneity amongst the processors **610**, **615**. For at least one embodiment, the various processors **610**, **615** may reside in the same die package.

Referring now to FIG. 7, shown is a block diagram of a second system **700** in accordance with an embodiment of the present invention. As shown in FIG. 7, multiprocessor system **700** is a point-to-point interconnect system, and includes a first processor **770** and a second processor **780** coupled via a point-to-point interface **750**. Each of processors **770** and **780** may be some version of the processors or integrated circuits as previously described or as one or more of the processors **610**, **615**.

While shown with only two processors **770**, **780**, it is to be understood that the scope of the present invention is not so limited. In other embodiments, one or more additional processors may be present in a given processor.

Processors **770** and **780** are shown including integrated memory controller units **772** and **782**, respectively. Processor **770** also includes as part of its bus controller units point-to-point (P-P) interfaces **776** and **778**; similarly, second processor **780** includes P-P interfaces **786** and **788**. Processors **770**, **780** may exchange information via a point-to-point (P-P) interface **750** using P-P interface circuits **778**, **788**. As shown in FIG. 7, IMCs **772** and **782** couple the processors to respec-

tive memories, namely a memory 732 and a memory 734, which may be portions of main memory locally attached to the respective processors.

Processors 770, 780 may each exchange information with a chipset 790 via individual P-P interfaces 752, 754 using point to point interface circuits 776, 794, 786, 798. Chipset 790 may also exchange information with a high-performance graphics circuit 738 via a high-performance graphics interface 739.

A shared cache (not shown) may be included in either processor or outside of both processors, yet connected with the processors via P-P interconnect, such that either or both processors' local cache information may be stored in the shared cache if a processor is placed into a low power mode.

Chipset 790 may be coupled to a first bus 716 via an interface 796. In one embodiment, first bus 716 may be a Peripheral Component Interconnect (PCI) bus, or a bus such as a PCI Express bus or another third generation I/O interconnect bus, although the scope of the present invention is not so limited.

As shown in FIG. 7, various I/O devices 714 may be coupled to first bus 716, along with a bus bridge 718 which couples first bus 716 to a second bus 720. In one embodiment, second bus 720 may be a low pin count (LPC) bus. Various devices may be coupled to second bus 720 including, for example, a keyboard and/or mouse 722, communication devices 727 and a storage unit 728 such as a disk drive or other mass storage device which may include instructions/code and data 730, in one embodiment. Further, an audio I/O 724 may be coupled to second bus 720. Note that other architectures are possible. For example, instead of the point-to-point architecture of FIG. 7, a system may implement a multi-drop bus or other such architecture.

Referring now to FIG. 8, shown is a block diagram of a system 800 in accordance with an embodiment of the present invention. FIG. 8 illustrates that the processors 870, 880 may include integrated memory and I/O control logic ("CL") 872 and 882, respectively and intercommunicate with each other via point-to-point interconnect 850 between point-to-point (P-P) interfaces 878 and 888 respectively. Processors 870, 880 each communicate with chipset 890 via point-to-point interconnects 852 and 854 through the respective P-P interfaces 876 to 894 and 886 to 898 as shown. For at least one embodiment, the CL 872, 882 may include integrated memory controller units. CLs 872, 882 may include I/O control logic. As depicted, memories 832, 834 coupled to CLs 872, 882 and I/O devices 814 are also coupled to the control logic 872, 882. Legacy I/O devices 815 are coupled to the chipset 890 via interface 896.

FIG. 9 depicts a tablet computing device 901 and a hand-held smartphone 902 each having a circuitry integrated therein as described in accordance with the embodiments. As depicted, each of the tablet computing device 901 and the hand-held smartphone 902 include a touchscreen interface 903 and an integrated processor 904 in accordance with disclosed embodiments.

For example, in one embodiment, a system embodies a tablet computing device 901 or a hand-held smartphone 902, in which a display unit of the system includes a touchscreen interface 903 for the tablet or the smartphone and further in which memory and an integrated circuit operating as an integrated processor are incorporated into the tablet or smartphone, in which the integrated processor implements one or more of the embodiments described herein for cutting senior store latency using store prefetching. In one embodiment, the integrated circuit described above or the depicted integrated processor of the tablet or smartphone is an integrated silicon

processor functioning as a central processing unit for a tablet computing device or a smartphone.

For example, in accordance with one embodiment, a tablet computing device includes: a display unit; and an integrated circuit which may operate as, for example, an out-of-order processor, and further in which the integrated circuit includes: a data bus, a pipeline to receive a store instruction in which the integrated circuit is to process out-of-order instructions and enforce in-order requirements for a cache; an Address Generation Unit (AGU) to calculate a physical address for the store instruction; and a fill buffer to pre fetch a cache line based on the store instruction and the calculated physical address before the store instruction retires from the pipeline.

FIG. 10 is a block diagram 1000 of an embodiment of tablet computing device, a smartphone, or other mobile device in which touchscreen interface connectors are used. Processor 1010 performs the primary processing operations. Audio subsystem 1020 represents hardware (e.g., audio hardware and audio circuits) and software (e.g., drivers, codecs) components associated with providing audio functions to the computing device. In one embodiment, a user interacts with the tablet computing device or smartphone by providing audio commands that are received and processed by processor 1010.

Display subsystem 1030 represents hardware (e.g., display devices) and software (e.g., drivers) components that provide a visual and/or tactile display for a user to interact with the tablet computing device or smartphone. Display subsystem 1030 includes display interface 1032, which includes the particular screen or hardware device used to provide a display to a user. In one embodiment, display subsystem 1030 includes a touchscreen device that provides both output and input to a user.

I/O controller 1040 represents hardware devices and software components related to interaction with a user. I/O controller 1040 can operate to manage hardware that is part of audio subsystem 1020 and/or display subsystem 1030. Additionally, I/O controller 1040 illustrates a connection point for additional devices that connect to the tablet computing device or smartphone through which a user might interact. In one embodiment, I/O controller 1040 manages devices such as accelerometers, cameras, light sensors or other environmental sensors, or other hardware that can be included in the tablet computing device or smartphone. The input can be part of direct user interaction, as well as providing environmental input to the tablet computing device or smartphone.

In one embodiment, the tablet computing device or smartphone includes power management 1050 that manages battery power usage, charging of the battery, and features related to power saving operation. Memory subsystem 1060 includes memory devices for storing information in the tablet computing device or smartphone. Connectivity 1070 includes hardware devices (e.g., wireless and/or wired connectors and communication hardware) and software components (e.g., drivers, protocol stacks) to the tablet computing device or smartphone to communicate with external devices. Cellular connectivity 1072 may include, for example, wireless carriers such as GSM (global system for mobile communications), CDMA (code division multiple access), TDM (time division multiplexing), or other cellular service standards). Wireless connectivity 1074 may include, for example, activity that is not cellular, such as personal area networks (e.g., Bluetooth), local area networks (e.g., WiFi), and/or wide area networks (e.g., WiMax), or other wireless communication.

Peripheral connections 1080 include hardware interfaces and connectors, as well as software components (e.g., drivers, protocol stacks) to make peripheral connections as a periph-

15

eral device (“to” **1082**) to other computing devices, as well as have peripheral devices (“from” **1084**) connected to the tablet computing device or smartphone, including, for example, a “docking” connector to connect with other computing devices. Peripheral connections **1080** include common or standards-based connectors, such as a Universal Serial Bus (USB) connector, DisplayPort including MiniDisplayPort (MDP), High Definition Multimedia Interface (HDMI), Firewire, etc.

FIG. **11** shows a block diagram illustrating the development of IP cores according to one embodiment. Storage medium **1130** includes simulation software **1120** and/or hardware or software model **1110**. In one embodiment, the data representing the IP core design can be provided to the storage medium **1130** via memory **1140** (e.g., hard disk), wired connection (e.g., internet) **1150** or wireless connection **1160**. The IP core information generated by the simulation tool and model can then be transmitted to a fabrication facility **1165** where it can be fabricated by a 3rd party to perform at least one instruction in accordance with at least one embodiment.

In some embodiments, one or more instructions may correspond to a first type or architecture (e.g., x86) and be translated or emulated on a processor of a different type or architecture (e.g., ARM). An instruction, according to one embodiment, may therefore be performed on any processor or processor type, including ARM, x86, MIPS, a GPU, or other processor type or architecture.

FIG. **12** illustrates how an instruction of a first type is emulated by a processor of a different type, according to one embodiment. In FIG. **12**, program **1205** contains some instructions that may perform the same or substantially the same function as an instruction according to one embodiment. However the instructions of program **1205** may be of a type and/or format that is different or incompatible with processor **1215**, meaning the instructions of the type in program **1205** may not be able to execute natively by the processor **1215**. However, with the help of emulation logic, **1210**, the instructions of program **1205** are translated into instructions that are natively capable of being executed by the processor **1215**. In one embodiment, the emulation logic is embodied in hardware. In another embodiment, the emulation logic is embodied in a tangible, machine-readable medium containing software to translate instructions of the type in the program **1205** into the type natively executable by the processor **1215**. In other embodiments, emulation logic is a combination of fixed-function or programmable hardware and a program stored on a tangible, machine-readable medium. In one embodiment, the processor contains the emulation logic, whereas in other embodiments, the emulation logic exists outside of the processor and is provided by a third party. In one embodiment, the processor is capable of loading the emulation logic embodied in a tangible, machine-readable medium containing software by executing microcode or firmware contained in or associated with the processor.

FIG. **13** is a block diagram contrasting the use of a software instruction converter to convert binary instructions in a source instruction set to binary instructions in a target instruction set according to embodiments of the invention. In the illustrated embodiment, the instruction converter is a software instruction converter, although alternatively the instruction converter may be implemented in software, firmware, hardware, or various combinations thereof. FIG. **13** shows a program in a high level language **1302** may be compiled using an x86 compiler **1304** to generate x86 binary code **1306** that may be natively executed by a processor with at least one x86 instruction set core **1316**. The processor with at least one x86 instruction set core **1316** represents any processor that can perform

16

substantially the same functions as a Intel processor with at least one x86 instruction set core by compatibly executing or otherwise processing (1) a substantial portion of the instruction set of the Intel x86 instruction set core or (2) object code versions of applications or other software targeted to run on an Intel processor with at least one x86 instruction set core, in order to achieve substantially the same result as an Intel processor with at least one x86 instruction set core. The x86 compiler **1304** represents a compiler that is operable to generate x86 binary code **1306** (e.g., object code) that can, with or without additional linkage processing, be executed on the processor with at least one x86 instruction set core **1316**. Similarly, FIG. **13** shows the program in the high level language **1302** may be compiled using an alternative instruction set compiler **1308** to generate alternative instruction set binary code **1310** that may be natively executed by a processor without at least one x86 instruction set core **1314** (e.g., a processor with cores that execute the MIPS instruction set of MIPS Technologies of Sunnyvale, Calif. and/or that execute the ARM instruction set of ARM Holdings of Sunnyvale, Calif.). The instruction converter **1312** is used to convert the x86 binary code **1306** into code that may be natively executed by the processor without at least one x86 instruction set core **1314**. This converted code is not likely to be the same as the alternative instruction set binary code **1310** because an instruction converter capable of this is difficult to make; however, the converted code will accomplish the general operation and be made up of instructions from the alternative instruction set. Thus, the instruction converter **1312** represents software, firmware, hardware, or a combination thereof that, through emulation, simulation or any other process, allows a processor or other electronic device that does not have an x86 instruction set processor or core to execute the x86 binary code **1306**.

While the subject matter disclosed herein has been described by way of example and in terms of the specific embodiments, it is to be understood that the claimed embodiments are not limited to the explicitly enumerated embodiments disclosed. To the contrary, the disclosure is intended to cover various modifications and similar arrangements as would be apparent to those skilled in the art. Therefore, the scope of the appended claims should be accorded the broadest interpretation so as to encompass all such modifications and similar arrangements. It is to be understood that the above description is intended to be illustrative, and not restrictive. Many other embodiments will be apparent to those of skill in the art upon reading and understanding the above description. The scope of the disclosed subject matter is therefore to be determined in reference to the appended claims, along with the full scope of equivalents to which such claims are entitled.

What is claimed is:

1. A method comprising:

receiving an instruction, wherein the received instruction is a store instruction, at an out-of-order processor, wherein the out-of-order processor enforces in-order requirements for a cache;

performing address generation and translation for the store instruction to calculate a physical address of the memory to be accessed by the store instruction; and

in response to the received instruction being the store instruction, executing a pre-fetch operation for a cache line data based on the store instruction and for the same calculated physical address while the store instruction is pending and before the store instruction retires.

17

2. The method of claim 1:
wherein the out-of-order processor comprises a pipeline;
wherein receiving the store instruction at the out-of-order
processor comprises receiving the store instruction at
the pipeline.
3. The method of claim 2, wherein executing the pre-fetch
for the cache line before the store instruction retires com-
prises executing the pre-fetch for the cache line before the
store instruction retires from the pipeline.
4. The method of claim 2, wherein the store instruction is
subject to one of the following events:
retirement from the pipeline pursuant to an accurate
instruction prediction;
a pipeline flush event pursuant to an instruction mis-pre-
diction, wherein the pipeline flush event results in the
store instruction rendered obsolete.
5. The method of claim 1, wherein executing the pre-fetch
for the cache line comprises requesting write authority to the
cache line within the cache.
6. The method of claim 5:
wherein the cache comprises a Level 1 cache (L1 cache);
and
wherein the executing the pre-fetch for the cache line is a
cache miss at the L1 cache.
7. The method of claim 6, wherein executing the pre-fetch
for the cache line further comprises one or more of:
executing the pre-fetch against a higher level cache respon-
sive to the cache miss; or
executing the pre-fetch against a system non-cache random
access memory (RAM) communicably interfaced with
the out-of-order processor responsive to the cache miss.
8. The method of claim 1, further comprising:
loading the cache line into the cache with write authority
responsive to executing the pre-fetch for the cache line.
9. The method of claim 8, further comprising:
writing the cache line into the cache after the store instruc-
tion retires; and
deallocating a store buffer for the cache line.
10. The method of claim 1, further comprising:
identifying a pre-existing pre-fetch operation for the same
cache line; and
promoting the pre-existing pre-fetch operation to a senior
store.
11. The method of claim 10, wherein promoting the pre-
existing pre-fetch operation to the senior store comprises:
identifying a fill buffer allocated to bring the cache line
from a memory into the cache pursuant to the pre-exist-
ing pre-fetch operation;
modifying an attribute of the fill buffer from indicating a
pre-fetch operation to indicating a senior store opera-
tion; and
wherein promoting the pre-existing pre-fetch operation to
the senior store fulfills the pre-fetch for the cache line.
12. The method of claim 10, wherein promoting the pre-
existing pre-fetch operation to the senior store comprises:
determining the pre-fetch for the cache line has not com-
pleted;
identifying a fill buffer allocated to bring the cache line
from a memory into the cache pursuant to the pre-exist-
ing pre-fetch operation;
overwriting an attribute of the fill buffer from indicating the
pre-existing pre-fetch operation with an attribute indi-
cating a senior store; and
wherein promoting the pre-existing pre-fetch operation to
the senior store fulfills the pre-fetch for the cache line.
13. The method of claim 1, wherein the out-of-order pro-
cessor enforces the in-order requirements for a cache by:

18

- requiring that stores to the cache pursuant to store instruc-
tions received at the out-of-order processor are made
visible external to the out-of-order processor in strict
order, wherein the strict order is identical in sequence to
a program order being executed by the out-of-order pro-
cessor.
14. The method of claim 1, further comprising:
determining the pre-fetch for the cache line has not com-
pleted;
determining a Global Observation point has been reached
for the pre-fetch; and
hiding a Global Observation point from being visible exter-
nal to the out-of-order processor preventing access to the
cache line from entities external to the out-of-order pro-
cessor.
15. The method of claim 14, further comprising:
determining all previous senior stores are associated with a
corresponding Global Observation point; and
un-hiding the Global Observation point for the pre-fetch
allowing the pre-fetch to receive a Global Observation
point time-stamp subsequent to all previous stores to
ensure in-order requirements of the out-of-order proces-
sor.
16. An integrated circuit comprising:
a data bus;
a pipeline to receive an instruction, wherein the received
instruction is a store instruction, wherein the integrated
circuit to process out-of-order instructions and enforce
in-order requirements for a cache;
an Address Generation Unit (AGU) to calculate a physical
address for the store instruction; and
a fill buffer to, in response to the received instruction being
the store instruction, pre-fetch a cache line based on the
same store instruction and for the same calculated physi-
cal address while the store instruction is pending and
before the store instruction retires from the pipeline.
17. The integrated circuit of claim 16, wherein the inte-
grated circuit comprises a Central Processing Unit (CPU).
18. The integrated circuit of claim 16, wherein the CPU
comprises an out-of-order processor.
19. The integrated circuit of claim 16, wherein the inte-
grated circuit comprises a central processing unit for one of a
tablet computing device or a smartphone.
20. A system comprising:
a display unit; and
an integrated circuit, wherein the integrated circuit com-
prises:
a data bus,
an out-of-order pipeline to receive an instruction,
wherein the received instruction is a store instruction,
wherein the integrated circuit to process out-of-order
instructions and enforce in-order requirements for a
cache,
an Address Generation Unit (AGU) to calculate a physi-
cal address for the store instruction, and
a fill buffer to, in response to the received instruction
being the store instruction, pre-fetch a cache line
based on the same store instruction and for the same
calculated physical address while the store instruction
is pending and before the store instruction retires from
the pipeline, wherein stores to a cache pursuant to a
store instruction are made visible external to the out-
of-order pipeline in strict order, wherein the strict
order is identical in sequence to a program order being
executed by the out-of-order pipeline.

21. The system of claim 20:
wherein the system embodies a tablet or a smartphone;
wherein the display unit comprises a touchscreen interface
of the tablet or the smartphone; and
wherein the integrated circuit is incorporated into the tablet 5
or smartphone.

22. A tablet computing device comprising:
a display unit; and
an integrated circuit, wherein the integrated circuit com-
prises: 10
a data bus,
a pipeline to receive an instruction, wherein the received
instruction is a store instruction, wherein the inte-
grated circuit to process out-of-order instructions and
enforce in-order requirements for a cache, 15
an Address Generation Unit (AGU) to calculate a physical
address for the store instruction, and
a fill buffer to, in response to the received instruction being
the store instruction, pre-fetch a cache line based on the
same instruction and for the same calculated physical 20
address while the store instruction is pending and before
the store instruction retires from the pipeline.

* * * * *